

3.Unsupervised Learning

Unsupervised learning is a type of machine learning in which the algorithm is not provided with labeled data or a specific output to predict. Instead, the algorithm is tasked with finding patterns and structure in the data on its own. This makes it a particularly powerful tool for discovering hidden relationships in large datasets where the structure and characteristics of the data are not well understood.

In unsupervised learning, the algorithm is provided with a set of input data and must use statistical techniques to identify similarities, differences, and patterns within the data. Clustering and dimensionality reduction are common techniques used in unsupervised learning. Clustering involves grouping similar data points together based on their characteristics. This can be useful for discovering natural groupings or clusters within the data, which can provide insights into the underlying structure of the data. Dimensionality reduction involves transforming the data into a lower-dimensional space while preserving as much of the original information as possible. This can be useful for reducing the complexity of the data and making it easier to analyze or visualize.

Unsupervised learning has a wide range of applications in various fields, including data mining, image and speech recognition, natural language processing, and anomaly detection. Some examples include identifying customer segments in marketing, detecting fraud in finance, and finding hidden patterns in medical data.




Figure 3.1: Unsupervised learning

3.1 Clustering

Clustering is a type of unsupervised learning in machine learning where the algorithm groups similar data points together based on their characteristics. The goal is to identify natural groupings or clusters within the data that can provide insights into the underlying structure of the data. Clustering is used in a wide range of applications, including data mining, image and speech recognition, natural language processing, and anomaly detection. It can be used to identify customer segments in marketing, detect fraud in finance, and find hidden patterns in medical data.

There are several algorithms used in clustering, and the choice of algorithm depends on the specific problem and the characteristics of the data. Some of the most popular clustering algorithms are:

- ✓ *K-Means Clustering*: K-means clustering is a simple and efficient algorithm that partitions the data into k clusters, where k is a user-defined parameter. The algorithm works by randomly selecting k initial centroids and assigning each data point to the nearest centroid. The centroids are then updated by computing the mean of the data points in each cluster, and the process is repeated until the centroids converge.
- ✓ *Hierarchical Clustering*: Hierarchical clustering is a more complex algorithm that creates a hierarchy of clusters, with the most similar data points grouped together at the lowest level and the most dissimilar data points at the highest level. The algorithm can be either agglomerative, where each data point starts as a separate cluster and is then merged with other clusters, or divisive, where all data points start in a single cluster and are then recursively split into smaller clusters.
- ✓ *Density-Based Clustering*: Density-based clustering algorithms identify clusters based on the density of the data points. The algorithm identifies regions of high density as clusters and separates them from regions of low density. The most popular density-based clustering algorithm is DBSCAN (Density-Based Spatial

Clustering of Applications with Noise), which groups together data points that are close to each other and have a high density.

- ✓ *Fuzzy Clustering*: Fuzzy clustering assigns each data point a degree of membership to each cluster, rather than assigning it to a single cluster. The algorithm calculates the degree of membership based on the distance between the data point and the center of each cluster.

K-means clustering

K-means clustering is a popular clustering algorithm used in machine learning to partition a dataset into k clusters. In this algorithm, each data point is assigned to the cluster with the nearest mean or centroid. The goal is to minimize the sum of the squared distances between the data points and their assigned centroids. Conceptually, k-means clustering can be thought of as trying to group similar data points together in order to identify natural groupings or clusters within the data. The algorithm works by iteratively refining the assignment of data points to clusters until convergence.




Figure 3.2: K-means clustering

Algorithm:

1. Initialize k centroids randomly.
2. Assign each data point to the nearest centroid.

3. Recalculate the centroids based on the mean of the data points assigned to each cluster.
4. Repeat steps 2 and 3 until convergence (i.e., when the assignments of data points to clusters no longer change or the maximum number of iterations is reached).

Here's an example of k-means clustering in action:

Let's say we have a dataset of 100 data points and we want to cluster them into 3 groups. We first randomly initialize the centroids of each cluster. Let's say the initial centroids are as follows:

Centroid 1: (2, 4)

Centroid 2: (8, 6)

Centroid 3: (4, 8)

We then assign each data point to the nearest centroid based on the Euclidean distance between the data point and each centroid. After the first iteration, the assignments might look something like this:

Cluster 1: (1, 3), (3, 5), (2, 4), (1, 5)

Cluster 2: (7, 6), (9, 5), (8, 7), (7, 5), (8, 5)

Cluster 3: (3, 8), (5, 9), (4, 8), (6, 8), (4, 7), (5, 7), (6, 7)

We then recalculate the centroids based on the mean of the data points assigned to each cluster. After the second iteration, the new centroids might look something like this:

Centroid 1: (2.25, 5)

Centroid 2: (7.8, 6),

Centroid 3: (4.57, 7.57)

We repeat the assignment and centroid calculation steps until convergence is reached.

Here's the Python code for implementing k-means clustering:

```
from sklearn.cluster import KMeans  
import matplotlib.pyplot as plt  
import numpy as np
```

Generate random data points

```
X = np.random.rand(100, 2)
```

Specify number of clusters

```
k = 3
```

Initialize k-means algorithm

```
kmeans = KMeans(n_clusters=k)
```

Fit the algorithm to the data

```
kmeans.fit(X)
```

Get the cluster labels and centroids

```
labels = kmeans.labels_
```

```
centroids = kmeans.cluster_centers_
```

Visualize the clusters

```
plt.scatter(X[:,0], X[:,1], c=labels)
```

```
plt.scatter(centroids[:,0], centroids[:,1], marker='*', s=300, c='r')
```

```
plt.show()
```

In this example, we generate a random dataset of 100 data points with 2 dimensions. We specify the number of clusters as 3 and initialize the k-means algorithm. We fit the algorithm to the data, get the cluster labels and centroids, and then visualize the clusters. The resulting plot shows the data points colored according to their assigned cluster and the centroids marked with a red star.

K-means clustering is a versatile technique that can be applied to a wide range of applications. Here are some examples:

- ✓ Image segmentation: K-means clustering can be used to segment images into regions with similar characteristics, such as color or texture.
- ✓ Customer segmentation: K-means clustering can be used to segment customers based on their buying behavior or demographics, which can help businesses tailor their marketing strategies.

- ✓ Anomaly detection: K-means clustering can be used to identify outliers or anomalies in datasets, which can be useful in fraud detection or quality control.
- ✓ Document clustering: K-means clustering can be used to cluster documents based on their content, which can help with tasks such as information retrieval or topic modeling.
- ✓ Recommendation systems: K-means clustering can be used to group similar items or users together, which can be used to make recommendations for products or content.

Hierarchical Clustering

Hierarchical clustering is a clustering algorithm that aims to create a hierarchy of clusters in a dataset. It starts by considering each data point as a separate cluster and then progressively merges clusters based on their similarity or distance.

Algorithm:

The algorithm for hierarchical clustering can be summarized as follows:

- ✓ Start with each data point as a separate cluster.
- ✓ Compute the similarity or distance between all pairs of clusters.
- ✓ Merge the two closest clusters into a new cluster.
- ✓ Update the similarity or distance matrix.
- ✓ Repeat steps 2-4 until all data points belong to a single cluster or until a desired number of clusters is obtained.

There are two main approaches to hierarchical clustering:

- ✓ Agglomerative (Bottom-Up) Approach:
 - a. Start with each data point as a separate cluster.
 - b. Compute the distance/similarity matrix between all pairs of clusters.
 - c. Merge the two closest clusters into a new cluster.
 - d. Update the distance/similarity matrix.
 - e. Repeat the merging and updating steps until all data points belong to a single cluster.

✓ Divisive (Top-Down) Approach:

- Start with all data points in a single cluster.
- Compute the distance/similarity matrix between all pairs of data points.
- Select the cluster with the largest dissimilarity and split it into two clusters.
- Repeat the splitting step until each data point is in its own cluster.




Figure 3.3: Hierarchical Clustering

Here's an example of implementing agglomerative hierarchical clustering using the scikit-learn library in Python:

```
from sklearn.cluster import AgglomerativeClustering
import numpy as np
# Generate random data points
X = np.random.rand(10, 2)
# Perform hierarchical clustering
clustering = AgglomerativeClustering(n_clusters=2)
clustering.fit(X)
# Get cluster labels
labels = clustering.labels_
print("Cluster Labels:", labels)
```

Output:

Cluster Labels: [1 0 0 0 1 1 0 1 0 1]

The output shows the cluster labels assigned to each data point.

Applications

Hierarchical clustering has various applications in different domains. Here are a few examples:

- ✓ **Image Segmentation:** Hierarchical clustering can be used to segment images based on similarities in color, texture, or other features. It helps in identifying different regions or objects within an image.
- ✓ **Document Clustering:** Hierarchical clustering can group similar documents based on their content, allowing for document organization and information retrieval tasks.
- ✓ **Customer Segmentation:** Hierarchical clustering helps in segmenting customers based on their buying patterns, demographics, or behavior, enabling targeted marketing strategies.
- ✓ **Bioinformatics:** Hierarchical clustering is used to analyze gene expression data and identify patterns or clusters of genes with similar expression profiles.

These applications demonstrate the versatility of hierarchical clustering in uncovering structures and relationships in different types of data. It helps in understanding complex systems and making data-driven decisions.

Density-based clustering

Density-based clustering is a type of unsupervised machine learning algorithm that groups together data points that are in high-density regions, separated by low-density regions. This algorithm is useful for identifying clusters of arbitrary shapes and sizes, and is particularly effective for datasets with noise or outliers. The most popular density-based clustering algorithm is DBSCAN (Density-Based Spatial Clustering of Applications with Noise). The DBSCAN algorithm has two key parameters: epsilon,

which specifies the maximum distance between two points for them to be considered part of the same cluster, and `min_samples`, which specifies the minimum number of points required to form a dense region.

Here's a high-level overview of the DBSCAN algorithm:

1. Select an arbitrary point in the dataset that has not been visited and identify all points within a distance of `epsilon` from it. If there are at least `min_samples` points within this distance, create a new cluster and add these points to it. Mark all of these points as visited.
2. For each point in the cluster, find all points within a distance of `epsilon` from it that have not been visited. If there are at least `min_samples` points within this distance, add them to the cluster and mark them as visited.
3. Repeat steps 1 and 2 until all points have been visited.
4. Any points that have not been visited are considered noise and are not part of any cluster.

The resulting clusters will be areas of high density in the dataset, separated by areas of low density.




Figure 3.4: Density-based clustering

Here's an example of density-based clustering using the DBSCAN algorithm in Python:

```
import numpy as np
import matplotlib.pyplot as plt
```

```

from sklearn.cluster import DBSCAN
# Generate some random data
np.random.seed(0)
X = np.vstack((np.random.randn(100, 2), 10 + np.random.randn(100, 2)))
# Compute DBSCAN
db = DBSCAN(eps=0.5, min_samples=5)
y_pred = db.fit_predict(X)
# Plot the results
plt.scatter(X[:, 0], X[:, 1], c=y_pred)
plt.title('DBSCAN Clustering')
plt.show()

```

This code generates 200 random data points in 2 dimensions, with 100 points in each of two separate Gaussian distributions. The DBSCAN algorithm is then applied with $\epsilon=0.5$ and $\text{min_samples}=5$. The resulting clusters are plotted in different colors:

In this example, the DBSCAN algorithm is able to identify the two distinct Gaussian distributions as separate clusters, even though they are close together and have overlapping points. The algorithm is also able to ignore the isolated points in the middle of the plot, which are considered noise.

Fuzzy clustering

Fuzzy clustering is a type of unsupervised machine learning algorithm that allows for soft clustering, meaning that each data point can belong to more than one cluster with a degree of membership between 0 and 1. This is different from hard clustering, such as K-means, where each data point is assigned to a single cluster. Fuzzy clustering is particularly useful when there is ambiguity in the data, or when a data point can belong to multiple clusters with different degrees of membership. This algorithm is commonly used in image segmentation, pattern recognition, and data compression.




Figure 3.5: Fuzzy clustering

The most popular fuzzy clustering algorithm is Fuzzy C-Means (FCM). The FCM algorithm assigns membership degrees to each data point for each cluster. The membership degrees are then used to calculate the centroid of each cluster, and the process is repeated until convergence.

Here's a high-level overview of the FCM algorithm:

1. Initialize the membership degrees for each data point randomly. The membership degrees must be between 0 and 1, and the sum of the membership degrees for each data point must equal 1.
2. Calculate the centroid of each cluster based on the membership degrees.
3. Update the membership degrees based on the distance between each data point and each cluster centroid.
4. Repeat steps 2 and 3 until convergence, where convergence is defined as a small change in the membership degrees or cluster centroids.
5. The final membership degrees can be used to assign each data point to one or more clusters.

Here's an example of fuzzy clustering using the FCM algorithm in Python:

```

import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
from sklearn.metrics import pairwise_distances_argmin_min

# Generate some random data
np.random.seed(0)
X, y = make_blobs(n_samples=200, centers=3, random_state=0)

# Apply FCM
def fcm(X, n_clusters, max_iter=100, m=2):
    # Initialize membership degrees randomly
    U = np.random.rand(len(X), n_clusters)
    U /= np.sum(U, axis=1)[:, np.newaxis]

    for i in range(max_iter):
        # Calculate cluster centroids based on membership degrees
        centroids = U.T.dot(X) / U.sum(axis=0)[:, np.newaxis]
        # Update membership degrees based on distance to cluster centroids
        distances = pairwise_distances_argmin_min(X, centroids)[1]
        U = 1.0 / (1.0 + (distances / np.max(distances) / m) ** 2)
        # Normalize membership degrees
        U /= np.sum(U, axis=1)[:, np.newaxis]
    # Return final membership degrees
    return U

U = fcm(X, n_clusters=3)

# Plot the results
plt.scatter(X[:, 0], X[:, 1], c=U)
plt.title('Fuzzy C-Means Clustering')

```

plt.show()

This code generates 200 random data points in 2 dimensions with 3 distinct clusters using the `make_blobs` function in `scikit-learn`. The FCM algorithm is then applied with `n_clusters=3`, `max_iter=100`, and `m=2`. The resulting clusters are plotted in different colors based on the membership degrees:

In this example, the FCM algorithm is able to assign each data point to multiple clusters with varying degrees of membership, allowing for more flexibility in the clustering.

3.2 Dimensionality Reduction




Figure 3.6: Dimensionality Reduction

Dimensionality reduction is a process of reducing the number of features or variables in a dataset while retaining its essence. The idea is to convert high-dimensional data into a lower-dimensional representation that still retains important information for modeling and analysis. Dimensionality reduction is essential for machine learning tasks because high-dimensional datasets often suffer from the "curse of dimensionality." This phenomenon can lead to overfitting, poor generalization, and increased computational complexity. There are two types of dimensionality reduction techniques:

- ❖ **Feature Selection:** It involves selecting a subset of the most important features from the original dataset based on statistical tests, domain knowledge, or other criteria.
- ❖ **Feature Extraction:** It involves transforming the original features into a new set of features that captures the essence of the data.

Feature selection

Feature selection is a type of dimensionality reduction technique that involves selecting a subset of the most important features from the original dataset while discarding the rest. The idea is to reduce the dimensionality of the dataset without losing relevant information.

Feature selection is important because high-dimensional datasets often suffer from the "curse of dimensionality," which can lead to overfitting, poor generalization, and increased computational complexity. By reducing the number of features, feature selection can improve the accuracy and efficiency of machine learning models.

There are several methods for feature selection, including:

Filter methods: These methods use statistical tests to rank the features based on their relevance to the target variable. Examples of filter methods include chi-squared test and correlation coefficient.

Filter methods are a category of feature selection techniques that evaluate each feature in a dataset independently based on some statistical measure, such as correlation or mutual information, to determine its relevance to the target variable. The goal of filter methods is to identify the subset of features that have the highest predictive power while minimizing redundancy and computational complexity.

The basic idea behind filter methods is to rank the features based on some metric that captures their relevance to the target variable. Commonly used metrics include:

1. Pearson correlation coefficient: measures the linear relationship between two variables.
2. Mutual information: measures the amount of information shared between two variables.
3. Chi-square test: measures the independence between two categorical variables.
4. ANOVA F-test: measures the difference in means between two or more groups.

Once the features are ranked, a threshold can be set to select the top N features, where N is a hyperparameter that can be tuned using cross-validation. The selected features can then be used as input to a machine learning algorithm to train a predictive model.

Filter methods are computationally efficient and can handle high-dimensional datasets with a large number of features. However, they do not take into account the interactions between features or the specific requirements of the learning algorithm being used. Therefore, filter methods may not always identify the optimal subset of features for a given problem, but they can provide a useful initial screening of the features and help reduce the dimensionality of the input space.

Wrapper methods: These methods use a machine learning algorithm to evaluate the importance of each feature by recursively training and testing the model on different subsets of the features.

Wrapper methods are a category of feature selection techniques that use a search algorithm, such as forward selection or backward elimination, to evaluate a subset of features based on their performance with a given machine learning algorithm. The goal of wrapper methods is to identify the subset of features that results in the highest predictive accuracy while accounting for interactions between features and the requirements of the learning algorithm.

The basic idea behind wrapper methods is to create a loop that iteratively adds or removes features based on their individual contribution to the predictive accuracy of the machine learning algorithm. The process typically involves the following steps:

- ❖ Initialize an empty set of selected features.
- ❖ Choose an initial subset of features to evaluate.
- ❖ Train a machine learning algorithm on the selected features and evaluate its performance using a cross-validation procedure.
- ❖ Use the performance metric to select the best subset of features and add them to the set of selected features.

- ❖ Repeat steps 3 and 4 with the new set of selected features until a stopping criterion is met.

Wrapper methods are computationally expensive and require a separate training and evaluation step for each subset of features. However, they can account for interactions between features and the specific requirements of the learning algorithm, which can result in a higher predictive accuracy than filter methods. Wrapper methods are typically used in situations where the number of features is small to moderate and the computational cost is not a major concern.

Embedded methods: These methods incorporate feature selection into the training process of the machine learning model. Examples of embedded methods include Lasso regression and decision tree-based feature selection.

Embedded methods are a category of feature selection techniques that incorporate the feature selection process into the training of a machine learning algorithm. Embedded methods use an optimization algorithm, such as gradient descent or Lasso regularization, to learn the best set of features for a given machine learning algorithm.

The basic idea behind embedded methods is to add a penalty term to the objective function of the machine learning algorithm that encourages the selection of a sparse set of features. The penalty term typically depends on a hyperparameter that controls the strength of the regularization and determines the degree of sparsity of the selected features.

Embedded methods can be used with a variety of machine learning algorithms, including linear and logistic regression, support vector machines, and neural networks. The specific regularization technique used depends on the learning algorithm and the nature of the problem being solved. For example, Lasso regularization is often used for linear regression and logistic regression, while elastic net regularization is often used for neural networks.

Embedded methods have several advantages over filter and wrapper methods. First, they can account for interactions between features and the specific requirements of the learning algorithm, which can result in a higher predictive accuracy. Second, they are computationally efficient because the feature selection process is integrated into the training of the machine learning algorithm. Finally, the selected features are often more robust to changes in the training data because they are learned as part of the optimization process.

However, embedded methods can be sensitive to the choice of hyperparameters, and the resulting feature set may not always be interpretable. Moreover, embedded methods may not be suitable for all types of machine learning algorithms or datasets, and the choice of regularization technique may depend on the nature of the problem being solved.

Feature selection can also be performed in conjunction with other dimensionality reduction techniques such as PCA and LDA. In this case, the most important features are selected from the reduced feature space. The general algorithm for feature selection can be summarized as follows:

1. Define the problem and the goal of feature selection.
2. Collect and preprocess the data.
3. Choose a feature selection method that is suitable for the problem.
4. Apply the feature selection method to rank or select the most important features.
5. Train and test a machine learning model on the selected features.
6. Evaluate the performance of the model on the task at hand.

Feature selection is widely used in various fields such as image processing, bioinformatics, finance, and natural language processing. It helps in reducing computational complexity, improving the accuracy of machine learning models, and enabling faster data analysis.

Feature extraction

Feature extraction is a key technique used in dimensionality reduction, which is the process of reducing the number of features in a dataset while retaining as much of the relevant information as possible. Dimensionality reduction is used in machine learning and data analysis to improve the performance and efficiency of algorithms, reduce the risk of overfitting, and enhance the interpretability of the results.

There are two main approaches to dimensionality reduction: feature selection and feature extraction. Feature selection involves selecting a subset of the original features based on some criterion, such as their correlation with the target variable or their mutual information with other features. Feature extraction involves transforming the original features into a lower-dimensional representation that captures the essential characteristics of the data.

Feature extraction methods for dimensionality reduction typically involve projecting the original data onto a lower-dimensional subspace, while preserving as much of the variance in the data as possible. Some commonly used feature extraction methods for dimensionality reduction include:

Principal Component Analysis (PCA): PCA is a linear transformation technique that projects the data onto a lower-dimensional subspace while maximizing the variance of the projected data. PCA identifies a set of orthogonal basis vectors, or principal components, that capture the directions of highest variance in the data.

Principal Component Analysis (PCA) is a widely used technique for feature extraction and dimensionality reduction. PCA is a linear transformation method that identifies the directions of highest variance in the data and projects the data onto a lower-dimensional subspace spanned by these directions.

The main goal of PCA is to reduce the dimensionality of the data while retaining as much of the original information as possible. PCA achieves this by finding a new set of orthogonal basis vectors, called principal components, that capture the directions of

highest variance in the data. The first principal component corresponds to the direction of highest variance, the second principal component corresponds to the direction of second-highest variance, and so on. The PCA algorithm can be summarized in the following steps:

- ✓ Standardize the data: PCA requires that the data be centered around zero and have unit variance. Therefore, the first step is to standardize the data by subtracting the mean and dividing by the standard deviation.
- ✓ Compute the covariance matrix: The covariance matrix measures the linear relationship between pairs of variables in the data. PCA computes the covariance matrix of the standardized data.
- ✓ Compute the eigenvectors and eigenvalues of the covariance matrix: The eigenvectors are the directions of highest variance in the data, and the eigenvalues represent the amount of variance explained by each eigenvector.
- ✓ Sort the eigenvectors in descending order of eigenvalues: The eigenvectors are sorted in order of decreasing importance, so that the first principal component corresponds to the direction of highest variance.
- ✓ Project the data onto the principal components: The data is projected onto the lower-dimensional subspace spanned by the first k principal components, where k is the desired number of dimensions.

PCA can be used for a variety of applications, including data visualization, feature extraction, and noise reduction. PCA is a powerful tool for dimensionality reduction because it identifies the directions of highest variance in the data, which often correspond to the most informative features. By projecting the data onto a lower-dimensional subspace, PCA can reduce the computational complexity of machine learning algorithms and improve their performance.

Linear Discriminant Analysis (LDA): LDA is a supervised dimensionality reduction technique that projects the data onto a lower-dimensional subspace while maximizing the

separation between classes. LDA identifies a set of basis vectors that maximize the between-class scatter and minimize the within-class scatter.

Linear Discriminant Analysis (LDA) is a supervised feature extraction method that is commonly used for dimensionality reduction. LDA is similar to PCA in that it projects the data onto a lower-dimensional subspace while preserving as much of the original information as possible. However, LDA takes into account the class labels of the data and seeks to maximize the separation between classes.

The main goal of LDA is to find a set of basis vectors that maximizes the separation between classes while minimizing the within-class scatter. The LDA algorithm can be summarized in the following steps:

- ✓ Standardize the data: LDA requires that the data be centered around zero and have unit variance. Therefore, the first step is to standardize the data by subtracting the mean and dividing by the standard deviation.
- ✓ Compute the class means: LDA computes the mean of each class in the standardized data.
- ✓ Compute the between-class scatter matrix: The between-class scatter matrix measures the separation between classes in the data.
- ✓ Compute the within-class scatter matrix: The within-class scatter matrix measures the variance within each class.
- ✓ Compute the eigenvectors and eigenvalues of the matrix $S^{-1}B$: LDA finds the eigenvectors and eigenvalues of the matrix $S^{-1}B$, where S is the within-class scatter matrix and B is the between-class scatter matrix.
- ✓ Sort the eigenvectors in descending order of eigenvalues: The eigenvectors are sorted in order of decreasing importance, so that the first linear discriminant corresponds to the direction of highest separation between classes.

- ✓ Project the data onto the linear discriminants: The data is projected onto the lower-dimensional subspace spanned by the first k linear discriminants, where k is the desired number of dimensions.

LDA can be used for a variety of applications, including classification, feature extraction, and data visualization. LDA is particularly useful when there is a small number of classes and a large number of features, as it can reduce the dimensionality of the data while improving the separation between classes. By projecting the data onto a lower-dimensional subspace that captures the essential characteristics of the data, LDA can improve the performance and efficiency of machine learning algorithms.

Non-negative Matrix Factorization (NMF): NMF is a matrix factorization technique that decomposes the data into a lower-dimensional representation of non-negative basis vectors and non-negative coefficients. NMF can be used to extract features that are sparse and interpretable.

Non-negative Matrix Factorization (NMF) is a popular feature extraction method used for dimensionality reduction in machine learning. NMF is particularly useful for data that is non-negative and sparse, such as image data, text data, and audio data.

The main idea behind NMF is to factorize a non-negative matrix X into two non-negative matrices W and H , such that $X \approx WH$. The matrix X represents the original data, while the matrix W represents the features or basis vectors that capture the essential characteristics of the data, and the matrix H represents the coefficients that weight the contribution of each feature to the original data. The NMF algorithm can be summarized in the following steps:

1. Initialize the matrices W and H with random non-negative values.
2. Update the matrix H by minimizing the objective function $\|X - WH\|^2$ with respect to H , subject to the constraint that H is non-negative.
3. Update the matrix W by minimizing the objective function $\|X - WH\|^2$ with respect to W , subject to the constraint that W is non-negative.

4. Repeat steps 2 and 3 until convergence.

NMF can be used for a variety of applications, including image processing, text mining, and signal processing. NMF is particularly useful for feature extraction because it can identify the underlying basis vectors that capture the essential characteristics of the data, while discarding the noise and irrelevant information.

One advantage of NMF over other feature extraction methods is that it results in a sparse representation of the data, where most of the coefficients in the matrix H are zero. This makes the representation more interpretable and reduces the dimensionality of the data. NMF is also useful for data compression and denoising, as it can extract the most informative features from the data while discarding the noise and redundant information.

Autoencoders: Autoencoders are a type of neural network that can be used for unsupervised dimensionality reduction. Autoencoders learn a compressed representation of the input data by encoding it into a lower-dimensional space and then decoding it back to the original space. The compressed representation can be used as a set of features for downstream machine learning algorithms.

Autoencoders are a type of neural network that can be used for feature extraction and dimensionality reduction. An autoencoder consists of two main components: an encoder and a decoder. The encoder maps the input data to a lower-dimensional representation, while the decoder maps the lower-dimensional representation back to the original input data.

The main idea behind autoencoders is to learn a compressed representation of the input data that captures the essential characteristics of the data while discarding the noise and irrelevant information. The autoencoder is trained to minimize the reconstruction error between the original input data and the reconstructed data, which encourages the autoencoder to learn a meaningful representation of the data. The autoencoder algorithm can be summarized in the following steps:

1. Define the architecture of the autoencoder: The autoencoder consists of an encoder network and a decoder network, which are typically symmetric. The encoder network maps the input data to a lower-dimensional representation, while the decoder network maps the lower-dimensional representation back to the original input data.
2. Train the autoencoder: The autoencoder is trained on a dataset of input data, by minimizing the reconstruction error between the original input data and the reconstructed data. The weights of the encoder and decoder networks are updated using backpropagation and gradient descent.
3. Use the encoder network for feature extraction: Once the autoencoder is trained, the encoder network can be used to extract features from new data. The lower-dimensional representation obtained by the encoder network can be used as a compressed representation of the data, which captures the essential characteristics of the data while discarding the noise and irrelevant information.

Autoencoders can be used for a variety of applications, including image processing, speech recognition, and natural language processing. Autoencoders are particularly useful for feature extraction and dimensionality reduction, as they can learn a compressed representation of the data that captures the essential characteristics of the data while reducing the dimensionality of the data. By extracting informative features from the data, autoencoders can improve the performance and efficiency of machine learning algorithms.

The choice of feature extraction method depends on the nature of the problem being solved, the type of data being analyzed, and the requirements of the machine learning algorithm. Feature extraction can improve the efficiency and accuracy of machine learning algorithms by reducing the dimensionality of the input data and by providing a more informative and representative set of features.

The general algorithm for dimensionality reduction techniques can be summarized as follows:

1. Define the problem and the goal of dimensionality reduction.
2. Collect and preprocess the data.
3. Choose a dimensionality reduction technique that is suitable for the problem.
4. Transform the data into a lower-dimensional representation.
5. Evaluate the performance of the technique on the task at hand.

Dimensionality reduction is widely used in various fields such as image processing, bioinformatics, finance, and natural language processing. It helps in reducing computational complexity, improving the accuracy of machine learning models, and enabling faster data analysis.

3.3 Association Rule Learning

Association Rule Learning is a type of unsupervised learning used to discover interesting relationships or associations between variables in large datasets.




Figure 3.7: Association Rule Learning Icon

Association Rule Learning is based on two main concepts: Support and Confidence.

Support: Support measures the frequency of occurrence of an itemset in the dataset. It is defined as the proportion of transactions in the dataset that contain the itemset. For example, if there are 100 transactions in the dataset, and 50 of them contain the itemset $\{A, B\}$, then the support of the itemset $\{A, B\}$ is $50/100 = 0.5$.

Confidence: Confidence measures the strength of the relationship between two itemsets. It is defined as the proportion of transactions containing the antecedent that also contain the consequent. For example, if there are 100 transactions in the dataset, and 50 of them contain the itemset $\{A, B\}$, and 40 of them also contain the itemset $\{A\}$, then the confidence of the association rule $\{A\} \rightarrow \{B\}$ is $40/50 = 0.8$.

Diagrams:

Association Rule Learning is often represented using a diagram called a lattice diagram or a Hasse diagram. This diagram represents the relationships between itemsets of different sizes. The diagram is constructed by drawing a node for each itemset, and connecting the nodes that are subsets of each other. The diagram allows us to visualize the hierarchy of itemsets and the relationships between them.

Algorithms:

Association Rule Learning is typically performed using the Apriori algorithm, which is an iterative algorithm that generates frequent itemsets and association rules. The algorithm works as follows:

1. Initialize the candidate itemsets: The algorithm starts by generating all possible itemsets of size 1, and calculating their support.
2. Prune infrequent itemsets: The algorithm removes any itemsets that do not meet the minimum support threshold.
3. Generate candidate itemsets of size $k+1$: The algorithm generates candidate itemsets of size $k+1$ by combining frequent itemsets of size k .
4. Prune infrequent candidate itemsets: The algorithm removes any candidate itemsets that do not meet the minimum support threshold.
5. Repeat steps 3-4 until no more frequent itemsets can be generated.
6. Generate association rules: The algorithm generates association rules for each frequent itemset, by considering all possible subsets of the itemset.

7. Evaluate the association rules: The algorithm evaluates the association rules based on their support and confidence, and selects the interesting rules based on the user-defined thresholds.

The Apriori algorithm is computationally intensive, as it involves generating and evaluating a large number of candidate itemsets. To improve the efficiency of the algorithm, several variations have been proposed, such as the FP-growth algorithm and the Eclat algorithm.

FOR AUTHOR USE ONLY